

# Experiment 4

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# Load the dataset
df = pd.read_csv(r'C:\Users\hp\OneDrive\Desktop\housing.csv') # Update with your file path

# Feature and target selection
X = df[['area']].values # Using 'area' as the feature
y = (df['price'] > df['price'].median()).astype(int).values # Binary target based on median price

# Normalize feature
X = (X - np.mean(X, axis=0)) / np.std(X, axis=0)

# Add a column of ones to include the intercept (bias) in the model
X_b = np.c_[np.ones((X.shape[0], 1)), X]

# Sigmoid function with clipping
def sigmoid(z):
    return 1 / (1 + np.exp(-np.clip(z, -250, 250))) # Clipping to avoid overflow

# Cost function with clipping
def cost_function(y, y_pred):
    y_pred = np.clip(y_pred, 1e-10, 1 - 1e-10) # Clipping to avoid log(0)
    m = len(y)
    return (-1/m) * np.sum(y * np.log(y_pred) + (1 - y) * np.log(1 - y_pred))

# Gradient Descent
def gradient_descent(X, y, learning_rate, num_iterations):
    m, n = X.shape
    theta = np.zeros(n)
    costs = []

    print(f'{"Iteration":<10}{"Theta0":<15}{"Theta1":<15}{"Cost Function"}')
    print("-" * 45)

    for i in range(num_iterations):
        predictions = sigmoid(X.dot(theta))
        errors = predictions - y
        gradients = (1/m) * X.T.dot(errors)
        theta -= learning_rate * gradients
        cost = cost_function(y, predictions)
        costs.append(cost)

        if i % 100 == 0: # Print every 100 iterations
            print(f"{i:<10}{theta[0]:<15.4f}{theta[1]:<15.4f}{cost:<20.4f}")

    return theta, costs

# Parameters
learning_rate = 0.01
num_iterations = 1000

# Run Gradient Descent
theta, costs = gradient_descent(X_b, y, learning_rate, num_iterations)

print(f"\nOptimal parameters (theta): {theta}")

# Plotting cost function over iterations
plt.figure(figsize=(10, 6))
plt.plot(range(num_iterations), costs, color='blue')
plt.xlabel('Number of Iterations')
plt.ylabel('Cost Function')
plt.title('Cost Function over Iterations')
plt.show()

# Plotting decision boundary
plt.figure(figsize=(10, 6))
plt.scatter(X, y, c=y, cmap='coolwarm', edgecolors='k', s=50)
x_values = np.linspace(X.min(), X.max(), 100)
y_values = (-1 / theta[1]) * (theta[0] + theta[1] * x_values)
plt.plot(x_values, y_values, label='Decision Boundary', color='red')
plt.xlabel('Area (normalized)')
plt.ylabel('Price (binary)')
plt.title('Logistic Regression Decision Boundary')
```

```
plt.legend()
plt.show()
```

Iteration	Theta0	Theta1	Cost Function
0	-0.0000	0.0024	0.6931
100	-0.0040	0.2142	0.6479
200	-0.0066	0.3816	0.6196
300	-0.0075	0.5170	0.6012
400	-0.0068	0.6287	0.5886
500	-0.0048	0.7225	0.5798
600	-0.0017	0.8023	0.5734
700	0.0020	0.8709	0.5686
800	0.0063	0.9304	0.5651
900	0.0108	0.9823	0.5623

Optimal parameters (theta): [0.01547946 1.02751984]



