

EXPERIMENT 1

```
In [19]: import pandas as pd

# Load the dataset using a raw string
df = pd.read_csv(r'C:\Users\hp\OneDrive\Desktop\housing.csv')
print("First 5 rows of the dataset:")
print(df.head())
```

First 5 rows of the dataset:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	\
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	

	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	no	yes	2	yes	furnished
1	no	yes	3	no	furnished
2	no	no	2	yes	semi-furnished
3	no	yes	3	yes	furnished
4	no	yes	2	no	furnished

IMPUTATION

```
In [20]: from sklearn.impute import SimpleImputer

# Separate numerical and categorical columns
numerical_cols = df.select_dtypes(include=['float64', 'int64']).columns
categorical_cols = df.select_dtypes(include=['object']).columns

# Impute missing values for numerical columns with mean
num_imputer = SimpleImputer(strategy='mean')
df[numerical_cols] = num_imputer.fit_transform(df[numerical_cols])

# Impute missing values for categorical columns with mode
cat_imputer = SimpleImputer(strategy='most_frequent')
df[categorical_cols] = cat_imputer.fit_transform(df[categorical_cols])

print("First 5 rows after imputation:")
print(df.head())
```

First 5 rows after imputation:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	\
0	13300000.0	7420.0	4.0	2.0	3.0	yes	no	
1	12250000.0	8960.0	4.0	4.0	4.0	yes	no	
2	12250000.0	9960.0	3.0	2.0	2.0	yes	no	
3	12215000.0	7500.0	4.0	2.0	2.0	yes	no	
4	11410000.0	7420.0	4.0	1.0	2.0	yes	yes	

	basement	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	no	no	yes	2.0	yes	furnished
1	no	no	yes	3.0	no	furnished
2	yes	no	no	2.0	yes	semi-furnished
3	yes	no	yes	3.0	yes	furnished
4	yes	no	yes	2.0	no	furnished

ANOMALY DETECTION

```
In [25]: from sklearn.ensemble import IsolationForest

# Verify the column name and use it in anomaly detection
print(df.columns) # Print column names to identify the correct one

# Choose a numerical column for anomaly detection, e.g., 'LotArea'
# Replace 'LotArea' with the correct column name if it differs
column_to_check = 'LotArea'

# Ensure the column exists in the DataFrame
if column_to_check in df.columns:
    clf = IsolationForest(contamination=0.05)
    df['anomaly'] = clf.fit_predict(df[[column_to_check]])
```

```
# Display rows marked as anomalies
anomalies = df[df['anomaly'] == -1]
print("Anomalous rows:")
print(anomalies.head())
else:
    print(f"Column '{column_to_check}' not found in DataFrame.")
```

```
Index(['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'mainroad',
      'guestroom', 'basement', 'hotwaterheating', 'airconditioning',
      'parking', 'prefarea', 'furnishingstatus'],
      dtype='object')
Column 'LotArea' not found in DataFrame.
```

STANDARDIZATION

In [22]: `from sklearn.preprocessing import StandardScaler`

```
# Standardize numerical columns
scaler = StandardScaler()
df[numerical_cols] = scaler.fit_transform(df[numerical_cols])

print("First 5 rows after standardization:")
print(df.head())
```

First 5 rows after standardization:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	\
0	4.566365	1.046726	1.403419	1.421812	1.378217	yes	no	
1	4.004484	1.757010	1.403419	5.405809	2.532024	yes	no	
2	4.004484	2.218232	0.047278	1.421812	0.224410	yes	no	
3	3.985755	1.083624	1.403419	1.421812	0.224410	yes	no	
4	3.554979	1.046726	1.403419	-0.570187	0.224410	yes	yes	

	basement	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	no	no	yes	1.517692	yes	furnished
1	no	no	yes	2.679409	no	furnished
2	yes	no	no	1.517692	yes	semi-furnished
3	yes	no	yes	2.679409	yes	furnished
4	yes	no	yes	1.517692	no	furnished

NORMALIZATION

In [23]: `from sklearn.preprocessing import MinMaxScaler`

```
# Normalize numerical columns
normalizer = MinMaxScaler()
df[numerical_cols] = normalizer.fit_transform(df[numerical_cols])

print("First 5 rows after normalization:")
print(df.head())
```

First 5 rows after normalization:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	\
0	1.000000	0.396564	0.6	0.333333	0.666667	yes	no	
1	0.909091	0.502405	0.6	1.000000	1.000000	yes	no	
2	0.909091	0.571134	0.4	0.333333	0.333333	yes	no	
3	0.906061	0.402062	0.6	0.333333	0.333333	yes	no	
4	0.836364	0.396564	0.6	0.000000	0.333333	yes	yes	

	basement	hotwaterheating	airconditioning	parking	prefarea	furnishingstatus
0	no	no	yes	0.666667	yes	furnished
1	no	no	yes	1.000000	no	furnished
2	yes	no	no	0.666667	yes	semi-furnished
3	yes	no	yes	1.000000	yes	furnished
4	yes	no	yes	0.666667	no	furnished

ENCODING

In [26]: `import pandas as pd`
`from sklearn.preprocessing import OneHotEncoder`

```
# Load the dataset
df = pd.read_csv(r'C:\Users\hp\OneDrive\Desktop\housing.csv') # Update the path if necessary

# Strip leading and trailing spaces from column names, if any
df.columns = df.columns.str.strip()

# Separate categorical columns
categorical_cols = df.select_dtypes(include=['object']).columns
```

```

# Initialize the OneHotEncoder with sparse_output=False
encoder = OneHotEncoder(sparse_output=False, drop='first') # Use sparse_output instead of sparse

# Fit and transform the categorical columns
encoded_features = encoder.fit_transform(df[categorical_cols])

# Create DataFrame for encoded features and concatenate with original DataFrame
encoded_df = pd.DataFrame(encoded_features, columns=encoder.get_feature_names_out(categorical_cols))
df_encoded = pd.concat([df.drop(categorical_cols, axis=1), encoded_df], axis=1)

print("First 5 rows after encoding:")
print(df_encoded.head())

```

First 5 rows after encoding:

	price	area	bedrooms	bathrooms	stories	parking	mainroad_yes	\
0	13300000	7420	4	2	3	2	1.0	
1	12250000	8960	4	4	4	3	1.0	
2	12250000	9960	3	2	2	2	1.0	
3	12215000	7500	4	2	2	3	1.0	
4	11410000	7420	4	1	2	2	1.0	

	guestroom_yes	basement_yes	hotwaterheating_yes	airconditioning_yes	\
0	0.0	0.0	0.0	1.0	
1	0.0	0.0	0.0	1.0	
2	0.0	1.0	0.0	0.0	
3	0.0	1.0	0.0	1.0	
4	1.0	1.0	0.0	1.0	

	prefarea_yes	furnishingstatus_semi-furnished	furnishingstatus_unfurnished
0	1.0	0.0	0.0
1	0.0	0.0	0.0
2	1.0	1.0	0.0
3	1.0	0.0	0.0
4	0.0	0.0	0.0

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js